



上海海事大学
SHANGHAI MARITIME UNIVERSITY

自然语言处理


2024–2025 学年第 2 学期



信息工程学院 谢雨波


PEFT 和 LoRA

PEFT (Parameter Efficient Fine-Tuning)



整体微调

更新模型的全部参数



PEFT

只更新模型的一小部分参数

PEFT：一类高效微调的方法

为什么只微调模型的一部分参数？

1. 对于大语言模型，微调全部参数不现实（成本太高）
 2. 如今的大语言模型大都是超参数化的（Over-parameterized）
- 只微调一部分参数也可以达到整体微调的效果

整体微调

- 假设有一个预训练因果语言模型 $P_\Phi(y|x)$
 - 例如基于 Transformer 解码器的 GPT 模型
- 对模型进行微调，应用于下游 NLP 任务
 - 需要训练数据集 $\{(x_i, y_i)\}_{i=1,\dots,N}$
- 在进行整体微调 (**Full Fine-tuning**) 时，使用梯度下降将 Φ_0 更新至 $\Phi_0 + \Delta\Phi$

$$\max_{\Phi} \sum_{(x,y)} \sum_{t=1}^{|y|} \log P_\Phi(y_t | x, y_{<t})$$

整体微调

- 对于每一个下游任务，都需要学习不同的 $\Delta\Phi$
- $|\Delta\Phi| = |\Phi_0|$
- 对于 GPT-3 来说， $|\Phi_0| = 1750$ 亿
 - 存储和部署模型的成本过高
- 如何进行优化？

LoRA

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li,
Shean Wang, Lu Wang, Weizhu Chen. [LoRA: Low-Rank Adaptation of Large Language Models](#). ICLR 2022.

- LoRA: **Low-Rank Adaptation**
- **核心思想**: 使用更少数目的参数 Θ 来表示 $\Delta\Phi$

$$\Delta\Phi = \Delta\Phi(\Theta), \quad |\Theta| \ll |\Phi_0|$$


- 因此, 求解 $\Delta\Phi$ 的任务可以转化为在 Θ 上进行寻优:

$$\max_{\Theta} \sum_{(x,y)} \sum_{t=1}^{|y|} \log P_{\Phi_0 + \Delta\Phi(\Theta)}(y_t | x, y_{<t})$$

LoRA

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen. [LoRA: Low-Rank Adaptation of Large Language Models](#). ICLR 2022.


- $W_0 \in \mathbb{R}^{d \times k}$: 预训练权重矩阵
- 使用一个低秩分解来限制 ΔW
 - $W_0 + \Delta W = W_0 + BA$
 - 其中 $B \in \mathbb{R}^{d \times r}, A \in \mathbb{R}^{r \times k}, r \ll \min(d, k)$
 - $h = W_0x + \Delta Wx = W_0x + BAx$
 - 对 ΔWx 进行缩放: $\frac{\alpha}{r}$ (α 和 r 均为超参数)
 - 只有 A 和 B 可训练, W_0 不参与训练



LoRA

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen. [LoRA: Low-Rank Adaptation of Large Language Models](#). ICLR 2022.

- 如果逐渐增加 LoRA 中的可训练参数，则训练 LoRA 将收敛至训练原模型
- 没有额外的推理延迟：从一个下游任务切换至另一个下游任务时，只需减去 BA 即可恢复 W_0 ，然后再加上一个不同的 $B'A'$
- 通常将 LoRA 应用于自注意力模块中的权重矩阵



LoRA

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li,
Shean Wang, Lu Wang, Weizhu Chen. [LoRA: Low-Rank Adaptation of Large Language Models](#). ICLR 2022.

- GPT-2 + LoRA 的实验结果

Model & Method	# Trainable Parameters	E2E NLG Challenge				
		BLEU	NIST	MET	ROUGE-L	CIDEr
GPT-2 M (FT)*	354.92M	68.2	8.62	46.2	71.0	2.47
GPT-2 M (Adapter ^L)*	0.37M	66.3	8.41	45.0	69.8	2.40
GPT-2 M (Adapter ^L)*	11.09M	68.9	8.71	46.1	71.3	2.47
GPT-2 M (Adapter ^H)	11.09M	67.3 _{±.6}	8.50 _{±.07}	46.0 _{±.2}	70.7 _{±.2}	2.44 _{±.01}
GPT-2 M (FT ^{Top2})*	25.19M	68.1	8.59	46.0	70.8	2.41
GPT-2 M (PreLayer)*	0.35M	69.7	8.81	46.1	71.4	2.49
GPT-2 M (LoRA)	0.35M	70.4 _{±.1}	8.85 _{±.02}	46.8 _{±.2}	71.8 _{±.1}	2.53 _{±.02}
GPT-2 L (FT)*	774.03M	68.5	8.78	46.0	69.9	2.45
GPT-2 L (Adapter ^L)	0.88M	69.1 _{±.1}	8.68 _{±.03}	46.3 _{±.0}	71.4 _{±.2}	2.49 _{±.0}
GPT-2 L (Adapter ^L)	23.00M	68.9 _{±.3}	8.70 _{±.04}	46.1 _{±.1}	71.3 _{±.2}	2.45 _{±.02}
GPT-2 L (PreLayer)*	0.77M	70.3	8.85	46.2	71.7	2.47
GPT-2 L (LoRA)	0.77M	70.4 _{±.1}	8.89 _{±.02}	46.8 _{±.2}	72.0 _{±.2}	2.47 _{±.02}

LoRA

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li,
Shean Wang, Lu Wang, Weizhu Chen. [LoRA: Low-Rank Adaptation of Large Language Models](#). ICLR 2022.

- GPT-3 + LoRA: 效果达到甚至超过整体微调

Model&Method	# Trainable Parameters	WikiSQL	MNLI-m	SAMSum
		Acc. (%)	Acc. (%)	R1/R2/RL
GPT-3 (FT)	175,255.8M	73.8	89.5	52.0/28.0/44.5
GPT-3 (BitFit)	14.2M	71.3	91.0	51.3/27.4/43.5
GPT-3 (PreEmbed)	3.2M	63.1	88.6	48.3/24.2/40.5
GPT-3 (PreLayer)	20.2M	70.1	89.5	50.8/27.3/43.5
GPT-3 (Adapter^H)	7.1M	71.9	89.8	53.0/28.9/44.8
GPT-3 (Adapter^H)	40.1M	73.2	91.5	53.2/29.0/45.1
GPT-3 (LoRA)	4.7M	73.4	91.7	53.8/29.8/45.9
GPT-3 (LoRA)	37.7M	74.0	91.6	53.4/29.2/45.1

LoRA

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li,
Shean Wang, Lu Wang, Weizhu Chen. [LoRA: Low-Rank Adaptation of Large Language Models](#). ICLR 2022.

- 将 LoRA 应用于 W_q 和 W_v 时，效果最好

		# of Trainable Parameters = 18M						
Weight Type	Rank r	W_q	W_k	W_v	W_o	W_q, W_k	W_q, W_v	W_q, W_k, W_v, W_o
WikiSQL ($\pm 0.5\%$)	8	70.4	70.0	73.0	73.2	71.4	73.7	73.7
MultiNLI ($\pm 0.1\%$)	8	91.0	90.8	91.0	91.3	91.3	91.3	91.7


- 即使 r 值很小，效果依旧可观

	Weight Type	$r = 1$	$r = 2$	$r = 4$	$r = 8$	$r = 64$
WikiSQL($\pm 0.5\%$)	W_q	68.8	69.6	70.5	70.4	70.0
	W_q, W_v	73.4	73.3	73.7	73.8	73.5
	W_q, W_k, W_v, W_o	74.1	73.7	74.0	74.0	73.9
MultiNLI ($\pm 0.1\%$)	W_q	90.7	90.9	91.1	90.7	90.7
	W_q, W_v	91.3	91.4	91.3	91.6	91.4
	W_q, W_k, W_v, W_o	91.2	91.7	91.7	91.5	91.4


优化器

参考：[李宏毅《机器学习》2021 春](#)


梯度下降



物理世界中的动量 (Momentum)



常规模梯度下降



初始位置: θ^0

计算梯度 g^0

移动到 $\theta^1 = \theta^0 - \eta g^0$


计算梯度 g^1

移动到 $\theta^2 = \theta^1 - \eta g^1$

⋮

梯度下降 + 动量

移动量：上一步的移动量 减去
当前的梯度



初始位置： θ^0

移动量 $m^0 = 0$

计算梯度 g^0

移动量 $m^1 = \lambda m^0 - \eta g^0$

移动到 $\theta^1 = \theta^0 + m^1$

计算梯度 g^1

移动量 $m^2 = \lambda m^1 - \eta g^1$

移动到 $\theta^2 = \theta^1 + m^2$

当前的移动量不仅依赖于梯度，还依赖于上一步的移动量

梯度下降 + 动量

移动量：上一步的移动量 减去
当前的梯度

m^i 是之前所有梯度的加权求和：

$$g^0, g^1, \dots, g^{i-1}$$

$$m^0 = 0$$

$$m^1 = -\eta g^0$$

$$m^2 = -\lambda \eta g^0 - \eta g^1$$

⋮

初始位置： θ^0

$$\text{移动量 } m^0 = 0$$

计算梯度 g^0

$$\text{移动量 } m^1 = \lambda m^0 - \eta g^0$$

$$\text{移动到 } \theta^1 = \theta^0 + m^1$$


计算梯度 g^1

$$\text{移动量 } m^2 = \lambda m^1 - \eta g^1$$

$$\text{移动到 } \theta^2 = \theta^1 + m^2$$

当前的移动量不仅依赖于梯度，还依赖于上一步的移动量



梯度下降 + 动量




没有临界点也很难训练？

损失函数为凸函数 →

学习率不能是一成不变的




$$\eta = 10^{-2}$$



$$\eta = 10^{-7}$$

不同的参数需要不同的学习率

对于某一个参数的梯度更新：



$$\theta_i^{t+1} \leftarrow \theta_i^t - \eta g_i^t$$

$$g_i^t = \frac{\partial L}{\partial \theta_i} \Big|_{\theta=\theta^t}$$

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$

依赖于参数

RMS (Root Mean Square) $\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$

$$\theta_i^1 \leftarrow \theta_i^0 - \frac{\eta}{\sigma_i^0} \mathbf{g}_i^0$$

$$\sigma_i^0 = \sqrt{(\mathbf{g}_i^0)^2} = |\mathbf{g}_i^0|$$

$$\theta_i^2 \leftarrow \theta_i^1 - \frac{\eta}{\sigma_i^1} \mathbf{g}_i^1$$

$$\sigma_i^1 = \sqrt{\frac{1}{2} [(\mathbf{g}_i^0)^2 + (\mathbf{g}_i^1)^2]}$$

$$\theta_i^3 \leftarrow \theta_i^2 - \frac{\eta}{\sigma_i^2} \mathbf{g}_i^2$$

$$\sigma_i^2 = \sqrt{\frac{1}{3} [(\mathbf{g}_i^0)^2 + (\mathbf{g}_i^1)^2 + (\mathbf{g}_i^2)^2]}$$

⋮

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$


$$\sigma_i^t = \sqrt{\frac{1}{t+1} \sum_{i=0}^t (g_i^t)^2}$$

RMS (Root Mean Square)

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$


$$\sigma_i^t = \sqrt{\frac{1}{t+1} \sum_{i=0}^t (g_i^t)^2}$$

在 AdaGrad 中使用



学习率根据参数动态地变化

实际情况中，损失函数非常复杂



RMSProp

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$

$$\theta_i^1 \leftarrow \theta_i^0 - \frac{\eta}{\sigma_i^0} \mathbf{g}_i^0$$

$$\sigma_i^0 = \sqrt{(\mathbf{g}_i^0)^2}$$

$0 < \alpha < 1$

$$\theta_i^2 \leftarrow \theta_i^1 - \frac{\eta}{\sigma_i^1} \mathbf{g}_i^1$$

$$\sigma_i^1 = \sqrt{\alpha (\sigma_i^0)^2 + (1 - \alpha) (\mathbf{g}_i^1)^2}$$

$$\theta_i^3 \leftarrow \theta_i^2 - \frac{\eta}{\sigma_i^2} \mathbf{g}_i^2$$

$$\sigma_i^2 = \sqrt{\alpha (\sigma_i^1)^2 + (1 - \alpha) (\mathbf{g}_i^2)^2}$$

⋮

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$

$$\sigma_i^t = \sqrt{\alpha (\sigma_i^{t-1})^2 + (1 - \alpha) (\mathbf{g}_i^t)^2}$$

RMSProp


$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$

$$\sigma_i^t = \sqrt{\alpha(\sigma_i^{t-1})^2 + (1-\alpha)(g_i^t)^2}$$

$g_i^1 \ g_i^2 \ \dots \ g_i^{t-1}$

$0 < \alpha < 1$

最近的梯度影响较大，过往的梯度影响较小



Adam: RMSProp + 动量

Diederik P. Kingma, Jimmy Ba. [Adam: A Method for Stochastic Optimization](#). ICLR 2015.

Algorithm 1: Adam, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1]$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

动量

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

RMSProp

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)



$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while


return θ_t (Resulting parameters)

适应性学习率 (Adaptive Learning Rate)



无适应性学习率


$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta}{\sigma_i^t} g_i^t$$
$$\sigma_i^t = \sqrt{\frac{1}{t+1} \sum_{i=0}^t (g_i^t)^2}$$



有适应性学习率



学习率规划 (Scheduling)

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta^t}{\sigma_i^t} g_i^t$$




学习率衰减 (Learning Rate Decay)

随着训练过程的推进，我们离目标越来越近，因此需要逐步降低学习率




学习率规划 (Scheduling)

$$\theta_i^{t+1} \leftarrow \theta_i^t - \frac{\eta^t}{\sigma_i^t} g_i^t$$



学习率衰减 (Learning Rate Decay)

随着训练过程的推进，我们离目标越来越近，因此需要逐步降低学习率



学习率预热 (Warm Up)

先增大学习率，再减小学习率

在初始阶段， σ_i^t 的估计 (Estimation) 有着较大的方差

AdamW

Ilya Loshchilov, Frank Hutter. [Decoupled Weight Decay Regularization](#). ICLR 2019.

- AdamW: Adam + Weight Decay (权重衰减)

模型泛化能力和收敛性能更好



Algorithm 2 Adam with L_2 regularization and Adam with decoupled weight decay (AdamW)

```
1: given  $\alpha = 0.001, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}, \lambda \in \mathbb{R}$ 
2: initialize time step  $t \leftarrow 0$ , parameter vector  $\theta_{t=0} \in \mathbb{R}^n$ , first moment vector  $m_{t=0} \leftarrow \theta$ , second moment
   vector  $v_{t=0} \leftarrow \theta$ , schedule multiplier  $\eta_{t=0} \in \mathbb{R}$ 
3: repeat
4:    $t \leftarrow t + 1$ 
5:    $\nabla f_t(\theta_{t-1}) \leftarrow \text{SelectBatch}(\theta_{t-1})$                                  $\triangleright$  select batch and return the corresponding gradient
6:    $\mathbf{g}_t \leftarrow \nabla f_t(\theta_{t-1}) + \lambda \theta_{t-1}$ 
7:    $\mathbf{m}_t \leftarrow \beta_1 \mathbf{m}_{t-1} + (1 - \beta_1) \mathbf{g}_t$                                  $\triangleright$  here and below all operations are element-wise
8:    $\mathbf{v}_t \leftarrow \beta_2 \mathbf{v}_{t-1} + (1 - \beta_2) \mathbf{g}_t^2$ 
9:    $\hat{\mathbf{m}}_t \leftarrow \mathbf{m}_t / (1 - \beta_1^t)$                                                $\triangleright \beta_1$  is taken to the power of  $t$ 
10:   $\hat{\mathbf{v}}_t \leftarrow \mathbf{v}_t / (1 - \beta_2^t)$                                                $\triangleright \beta_2$  is taken to the power of  $t$ 
11:   $\eta_t \leftarrow \text{SetScheduleMultiplier}(t)$                                  $\triangleright$  can be fixed, decay, or also be used for warm restarts
12:   $\theta_t \leftarrow \theta_{t-1} - \eta_t \left( \alpha \hat{\mathbf{m}}_t / (\sqrt{\hat{\mathbf{v}}_t} + \epsilon) + \lambda \theta_{t-1} \right)$ 
13: until stopping criterion is met
14: return optimized parameters  $\theta_t$ 
```
